

Event Ticketing Platform

Solution Architecture Assessment

February 2026

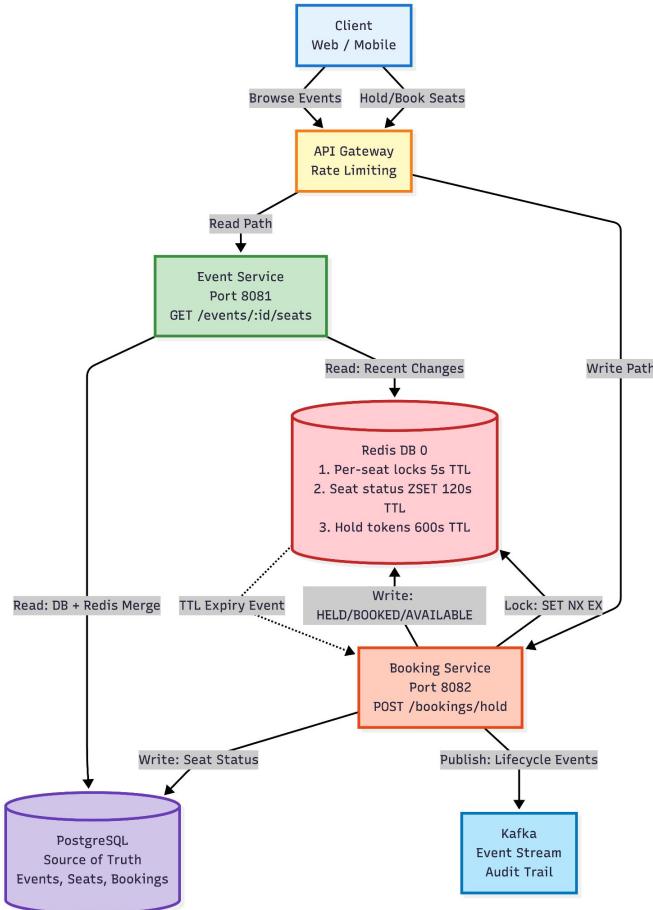
Problem & Scenarios

Category	Scenario	Strategic Rationale
Write Path	Seat Hold (e.g 10-Minute TTL)	Addresses high-concurrency locking, distributed coordination, expiry handling, and event-driven workflows
Read Path	Event Discovery with Caching	Optimizes search latency, availability reads, and cache consistency strategies
Challenge	Business Risk / Impact	
Double Booking	Revenue loss and reputational damage due to concurrent seat allocation conflicts	
Abandoned Holds	Inventory lock leading to reduced conversion and revenue leakage	
Lack of Audit Trail	Operational risk in dispute resolution and regulatory traceability	

Success Criteria (North Star Outcomes)

Zero double bookings | Automated seat release via deterministic TTL enforcement | End-to-end booking lifecycle auditability

Architectural Overview



Architecture Style

Microservices-based, event-driven system separating Event and Booking domains for scalability and fault isolation.

Concurrency Control

Redis distributed locks + ACID DB transactions ensure seat integrity under high contention.

Event-Driven Workflow

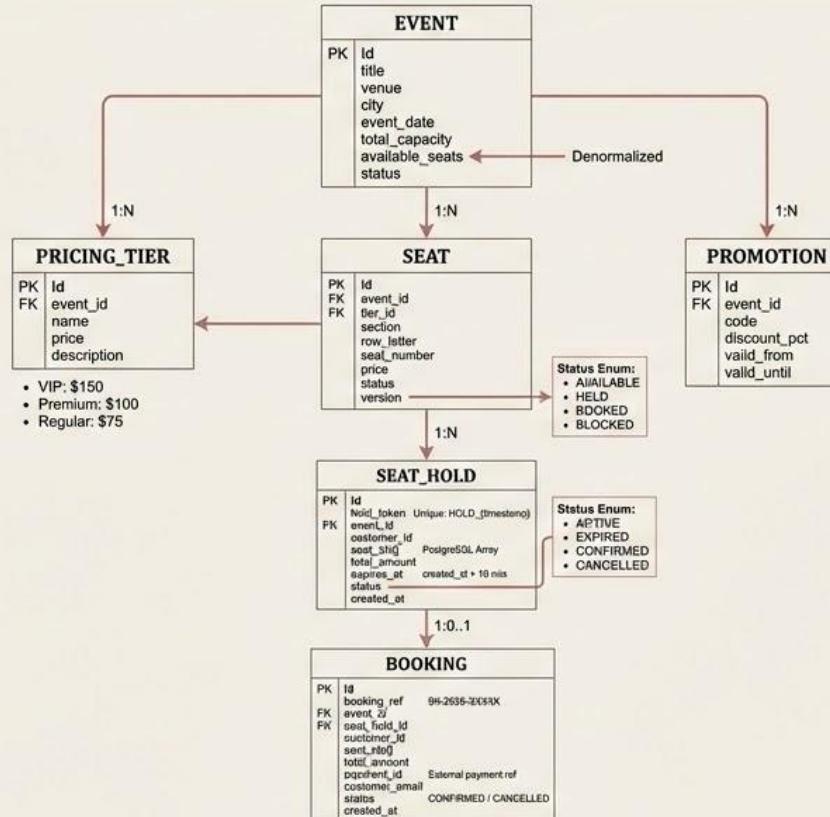
Kafka publishes lifecycle events:

SEAT-HOLD-CREATED, SEAT-HOLD-CONFIRMED, SEAT-HOLD-CANCELLED, SEAT-HOLD-EXPIRED, BOOKING-CONFIRMED

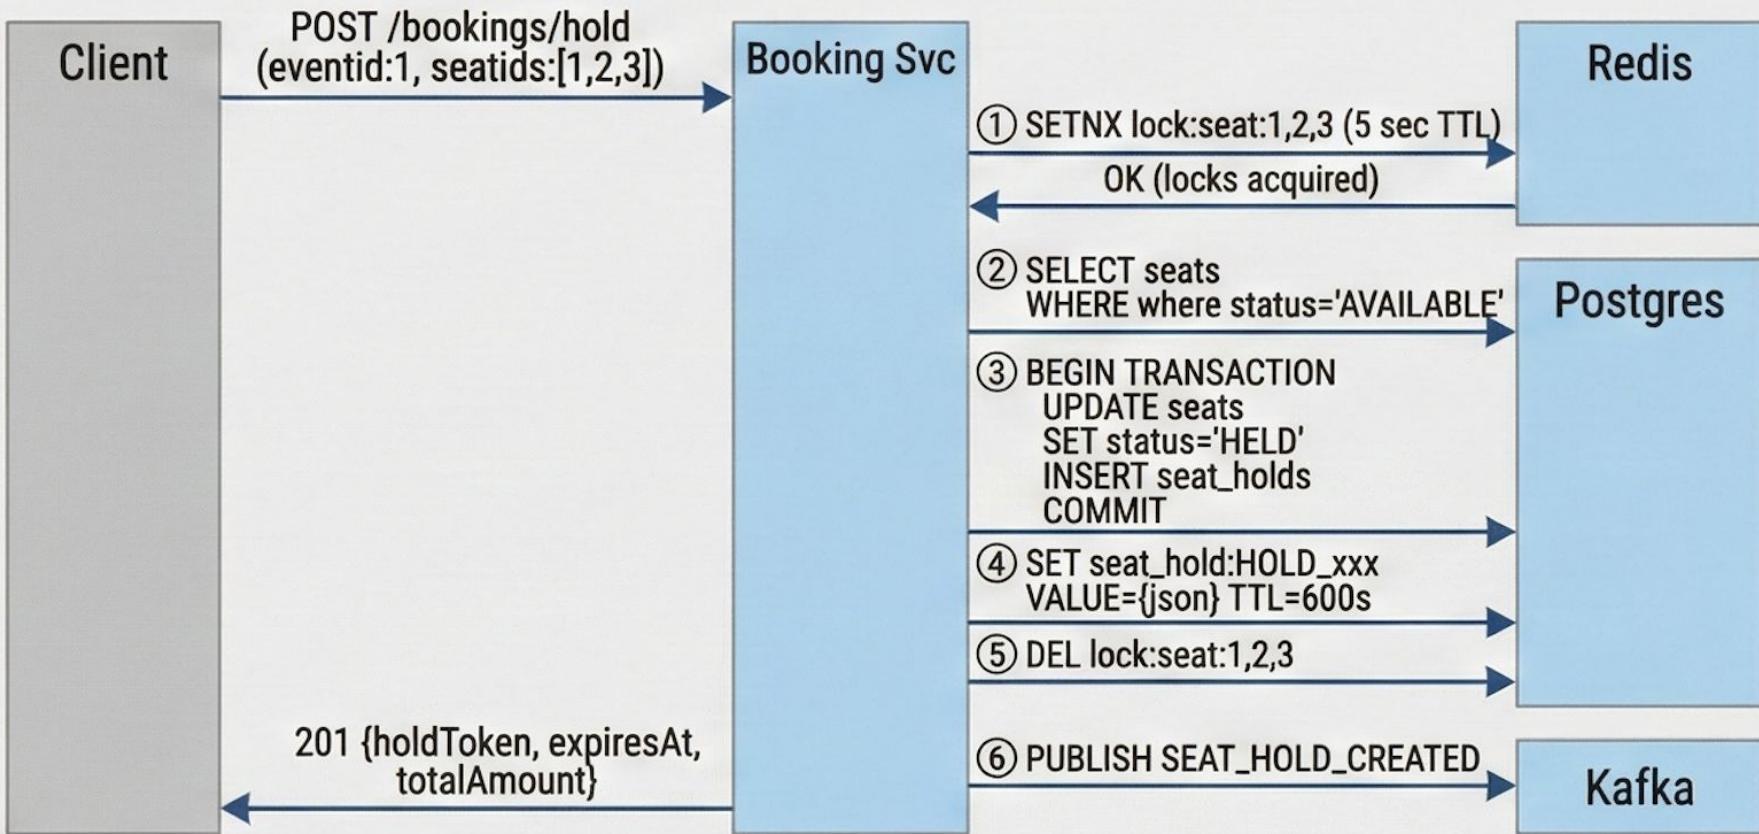
Resilience & Recovery

TTL-based auto release, idempotent APIs, and service isolation prevent deadlocks and duplicate bookings.

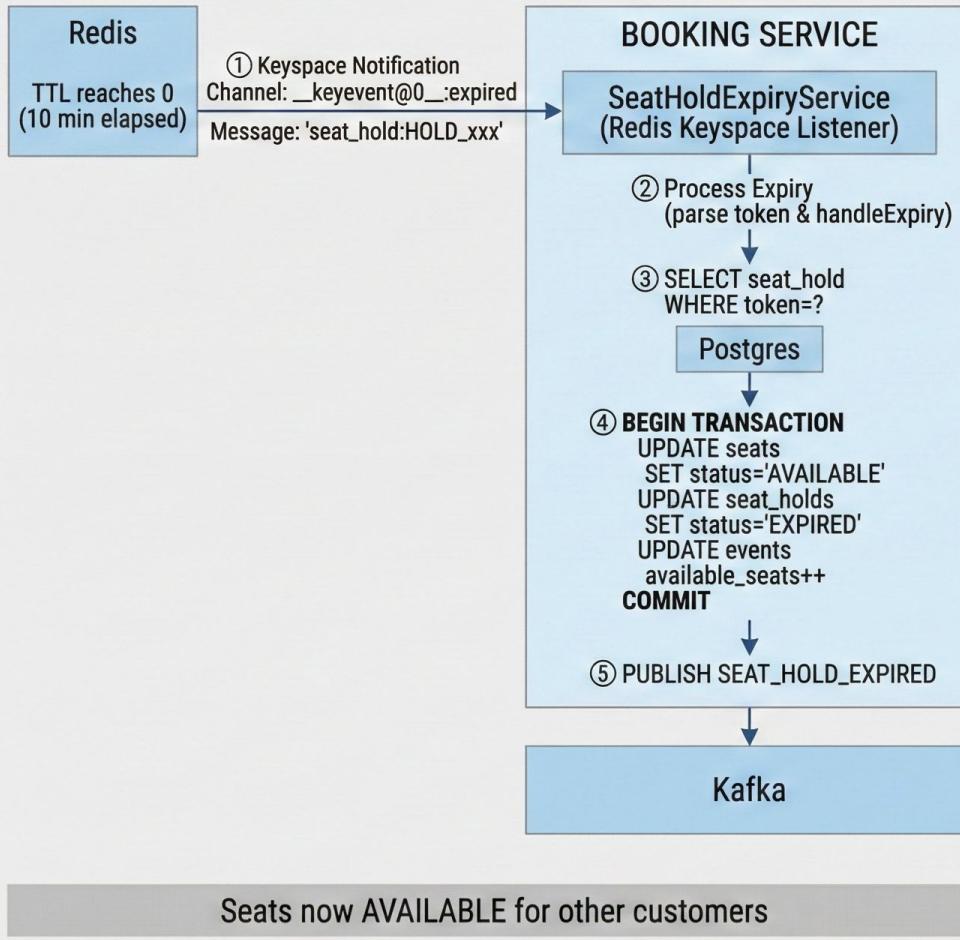
Data Entity Design



Scenario Flow - Seat Hold (Write)



Scenario Flow - Auto-Expiry (Timeout)

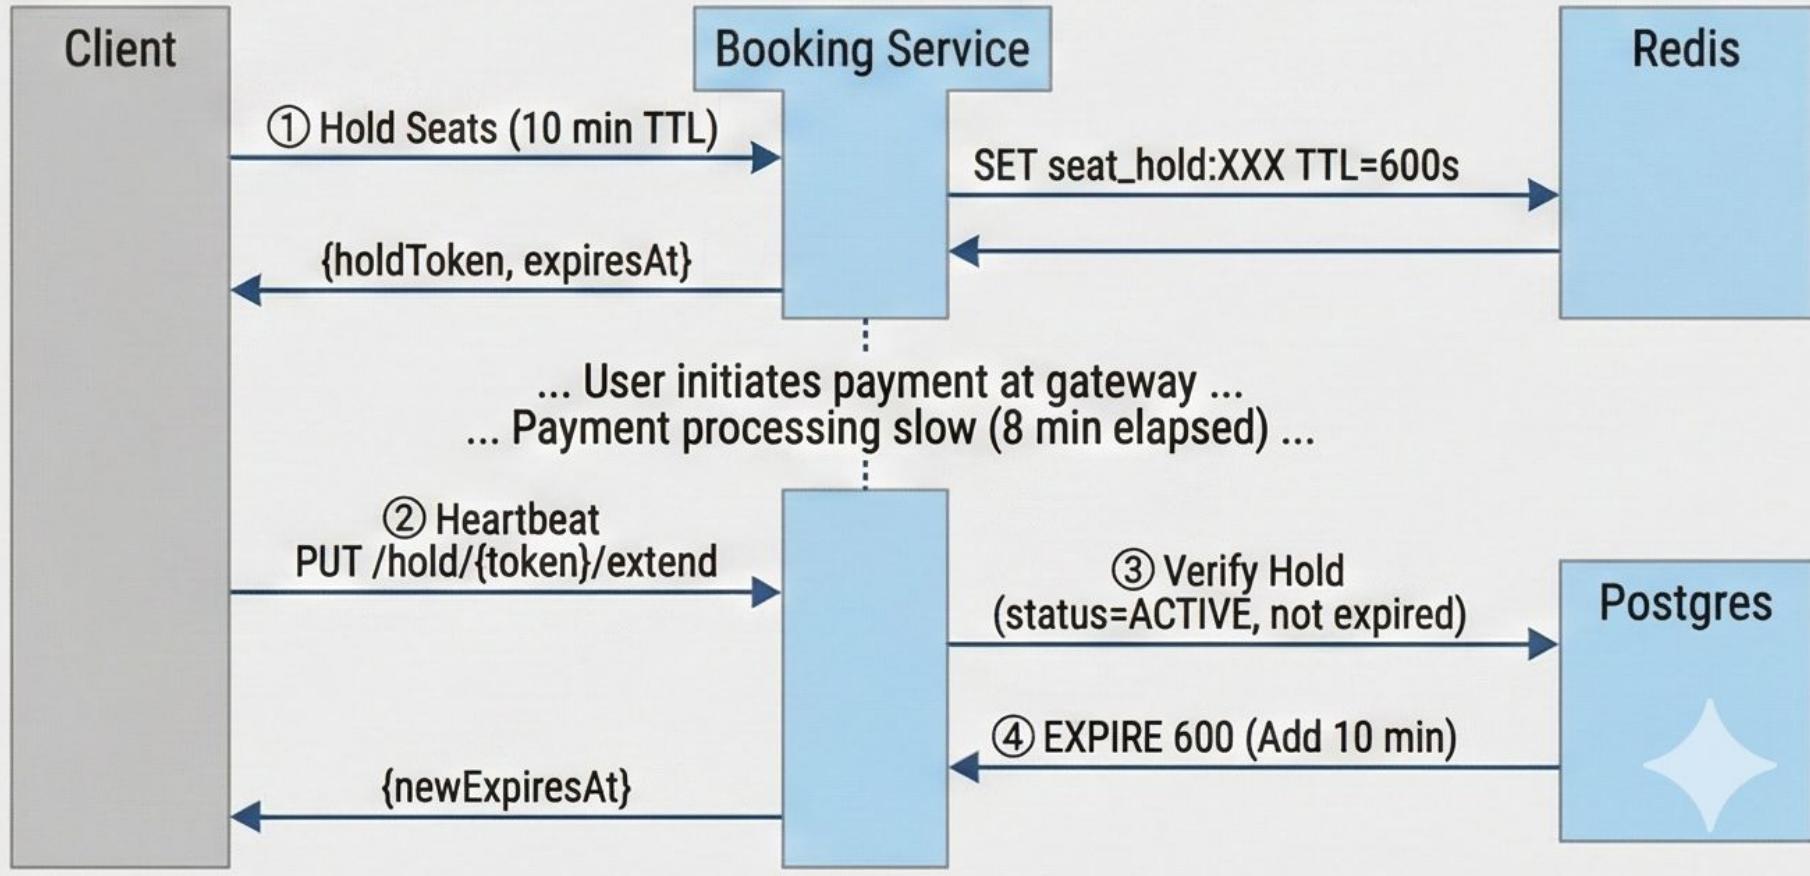


Assumption :

Booking not completed within Hold-Window

Redis key pattern: `seat:{eventId}:{seatId}:HELD`

Sequence 2: Heartbeat Extension for Slow Payment Gateways



Core Architectural Decision

Database as Source of Truth

All booking confirmations validated via **conditional DB update** Prevents double booking under race conditions Redis used only for coordination, not durability Strong consistency enforced at booking boundary

Deterministic Expiry & Self-Healing

Seat holds enforced with TTL Delayed expiry validation via background consumer Periodic reconciliation job (Redis ↔ DB drift correction) System remains consistent even after crashes

Distributed Locks

`SET NX EX` with unique lock token (UUID)
Lua-based safe unlock (ownership verification)
Short lock TTL (e.g., 30s) to avoid deadlocks
Designed to handle high-contention flash-sale scenarios

Idempotent by Design State Machine

The holdToken acts as a natural, single-use idempotency key for the confirm flow. A retry against an already-confirmed hold is rejected by the `isActive()` pre-check and the WHERE `status='HELD'` guard. domain state itself prevents duplicate effects

Key Architectural Trade-offs

PostgreSQL for Transactional Data

Strong ACID guarantees for bookings

Trade-off: Scaling requires partitioning

Redis for Distributed State

Fast locks, TTL-based holds, Real-Time data

Trade-off: Additional infrastructure complexity

Pessimistic Locking + Distributed Locks

Prevent double booking at multiple layers, supports Redis unavailability

Trade-off: Retry logic, potential contention

Microservices Architecture

Independent scaling, fault isolation

Trade-off: Increased operational complexity

Event-Driven Async Processing

Decouple notifications, analytics

Trade-off: Eventual consistency

API Design - Complete User Journey

GET /api/events/search

Browse events by city, date, category

200 OK with paginated results

GET /api/events/{id}/seats

View seat layout and availability

200 OK with seat map

POST /api/bookings/hold-seats

Hold selected seats (10 min TTL)

200 OK with hold ID

POST /api/bookings/{hold-token}/confirm

Create booking from hold

201 Created with booking ref

GET /api/bookings/{id}

Get booking details

200 OK with full details

POST /api/bookings/hold/{hold-token}?customerid={castID}

Cancel hold & return to AVAILABLE pool

200 OK with clear hold

Failure Scenario

1. Service Crash After Lock Acquisition

- Scenario: Lock acquired in Redis, but DB write fails/crashes.
- Risk: Seat locked in Redis, not in DB (state drift).

Mitigation:

1. Persist DB immediately after lock (fail-fast).
2. On DB failure → Safe unlock using Lua script:
if redis.call("GET", KEYS[1]) == ARGV[1] then
3. return redis.call("DEL", KEYS[1])
4. end
5. Reconciliation job (every 5 min):
 - Find Redis locks without DB records → Delete lock.
 - Find DB HELD records without Redis locks → Recreate/release.

2. Redis Crash / Restart

- Scenario: All locks lost after Redis restart.
- Risk: DB and Redis state drift (DB=HELD, no lock).

Mitigation:

1. DB is source of truth; Redis is derived/cache.
2. Periodic reconciliation job:
 - DB=HELD & no Redis lock → Recreate lock or mark expired.
 - Redis lock & DB=AVAILABLE → Delete stale lock.
3. Use Redis persistence (AOF/RDB) for faster recovery.
4. Fallback: Use DB row locks if Redis is down (degraded mode).

Failure Scenario continued..

3. Lock Expiry During Payment Processing

Scenario: The distributed lock (e.g., Redis lock) expires before the full payment transaction is finalized.

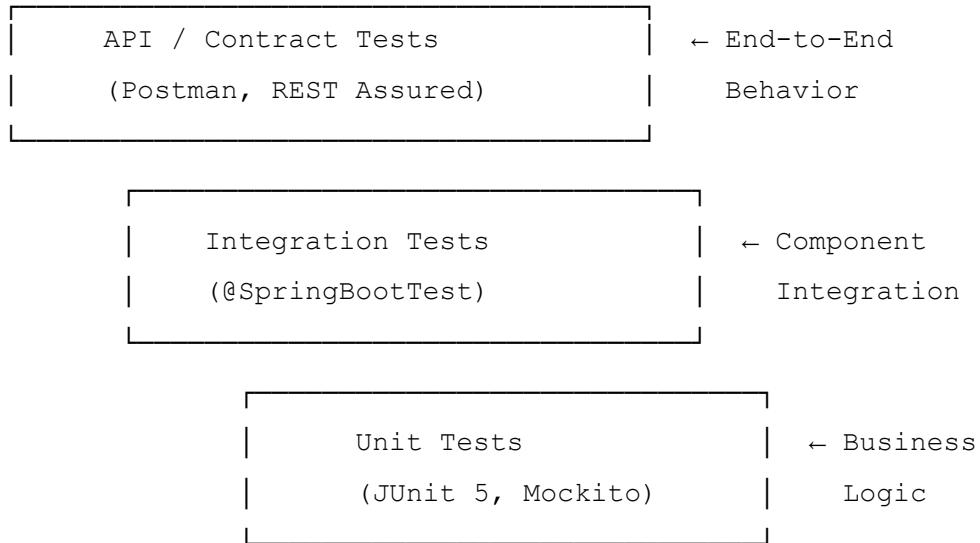
Mitigation:

- Enforce mandatory database conditional confirmation:-

```
UPDATE seats SET status='BOOKED' WHERE id IN (1,2,3) AND status='HELD'
```

- Optionally implement a Time-To-Live (TTL) heartbeat mechanism to extend the lock duration for protracted payment processes.
- **The booking attempt must be rejected if the conditional database update fails**

Testing Strategy and Quality Assurance



Thank You